

Uber Graph Benchmark

Statistically representative graph
generation and benchmarking

Chris Lu
Joshua Shinavier


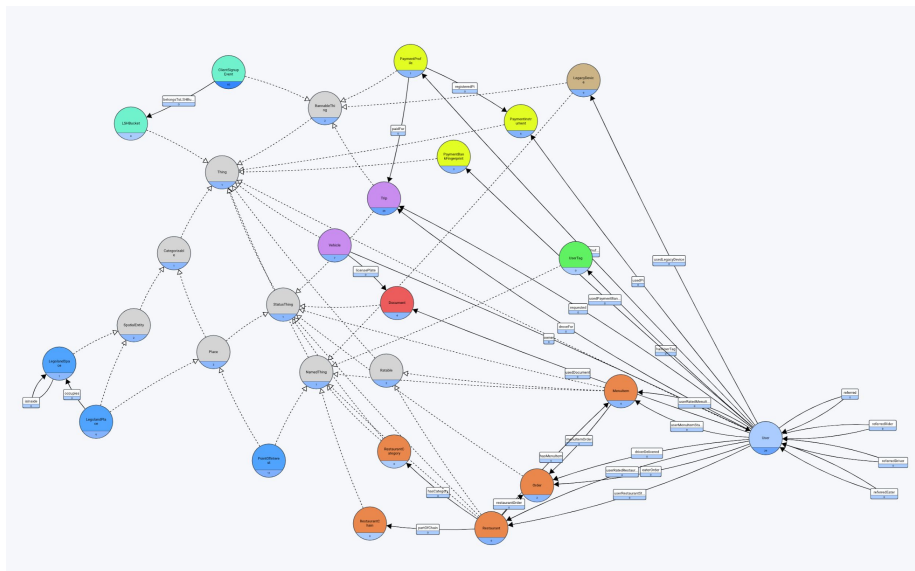


UBER

Agenda

- Why we need a new benchmarking framework
- How it works
- Invitation to participate

Let's put our data in a graph database!



Just pick one graph database?

Pick one graph database!

- Various graph database vendors.
- It's possible to build a graph database on key-value stores.
 - Various key-value stores.
 - Graph database on MySQL database
- We already have X key-value store, how does it compare to Y graph database?

It's complicated to choose one

- What is the read performance?
 - Data size: vertex/edge count, properties
 - Edge distribution
 - Hardware
 - Query pattern:
 - Concurrency
 - # hops
 - Fanout limit
 - Edge filtering
- What is the write performance?
- ...

The answer? It's complicated.

- What is the read performance?
- What is the write performance?
 - Graph Storage is usually denormalized
 - Denormalize data to optimize for performance
 - Storing adjacency list twice.
 - Build vertex-centric index.
 - Write amplification, transaction for consistency
- ...

More questions

- What is the read performance?
- What is the write performance?
- What is the read performance during heavy writes?
- What is the storage cost? Capacity planning? SKU?

Read the marketing materials

- Count keywords: "scalable", "fast", "in memory", "distributed"
- This is what Artificial Intelligence does.
- Hope your CTO does not do this...

Read other benchmarks, but

- Data is different
 - Schema
 - Edge distribution
 - Data size
- Hardware
 - Network
 - CPU
 - Memory
 - Disk
- Access patterns
 - ...

Which test to trust?

Test with your own data!

How about a DIY benchmark?

- * learn how to load data into the system
- * learn how to query data
- * build benchmark system {
 - * benchmark loading data
 - * benchmark querying data
 - * benchmark loading and querying at the same time}

How about a DIY benchmark?

for each system {

- * learn how to load data into the system

- * learn how to query data

- * build benchmark system {

 - * benchmark loading data

 - * benchmark querying data

 - * benchmark loading and querying at the same time

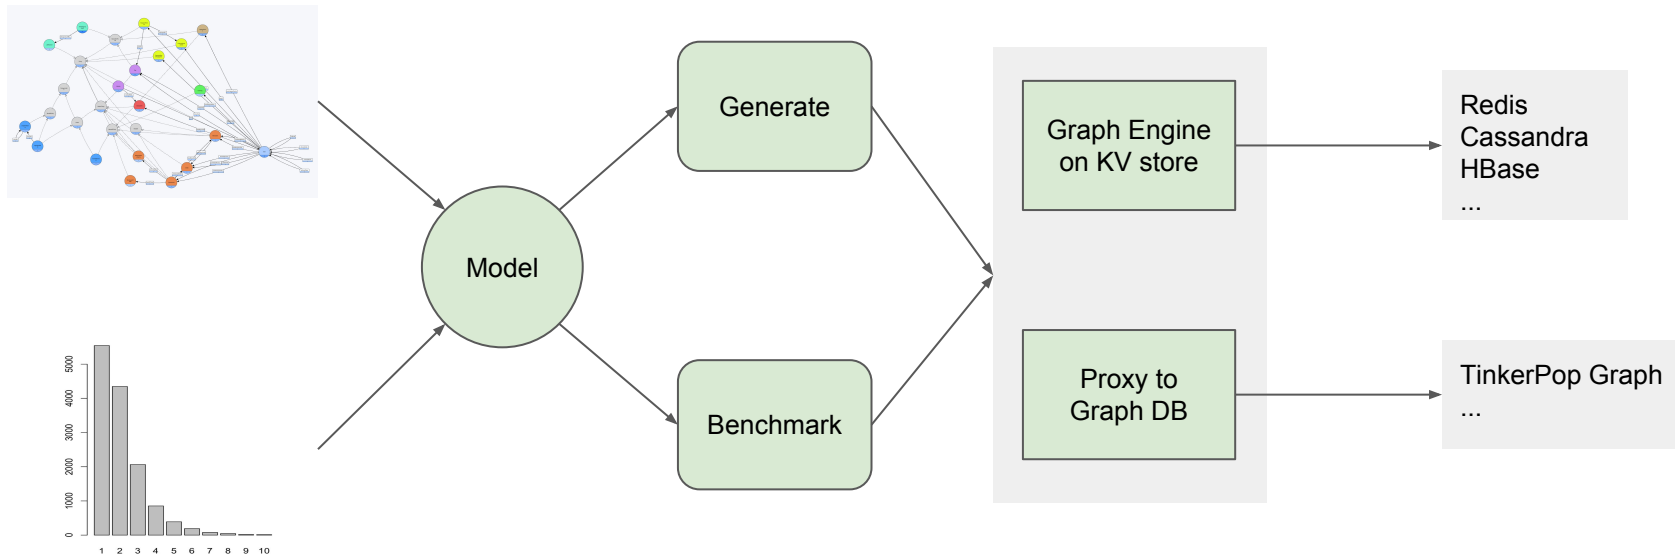
- }

}

We need a graph benchmark framework to

- Generate a graph representative of your data
 - Any size (as large as you need)
 - Any schema
 - Any distribution
 - Repeatable
- Store data into
 - Any graph store
 - Any key-value store
- Run k-hop subgraph on generated graph, starting from one random vertex

Uber Graph Benchmark Architecture



Graph Gen: Schema-Aware

- Create a statistical model of the dataset capturing:
 - Schema (vertex and edge labels, properties)
 - Label/key probabilities and frequency distributions
- Use the model to generate graphs of any size
- Generated graph should be “close enough” to real data to be used as a substitute

Graph Gen: Define Graph Schema

- Composable from different teams
- Example

entities:

- **label:** User

relations:

- **label:** isDriver

description: Whether a user is a driver

to: core.Boolean

- **label:** driverStatus

description: Current status of the driver (Active, Rejected, Waitlisted, etc.)

to: core.String

relations:

- **label:** requested

description: The relationship between a user and a trip he or she requested

extends:

- core.relatedTo

from: users.User

to: Trip

cardinality: OneToMany

Graph Gen: Define Graph Data Distribution

- Vertex

vertices:

- type: users.User
weight: 13
- type: trips.Trip
weight: 8
- type: documents.Document
weight: 13

Graph Gen: Define Graph Data Distribution

- Vertex

vertices:

- type: users.User
weight: 13
- type: trips.Trip
weight: 8
- type: documents.Document
weight: 13

- Property

properties:

- type: users.driverStatus
values:
 - value: "Active"
weight: 90
 - value: "Rejected"
weight: 5
 - value: "Waitlisted"
weight: 5

Generate Vertices

- Vertex id generation
 - Customizable for each graph db.
 - $Id = F(\text{String vertexLabel}, \text{long sequencedNumber})$
- Property types
 - Basic: Boolean, Date, Decimal, Double, Float, Long, String
 - Structured: Email, PhoneNumber, UnixTime, Year
 - List of values with weights.
- Vertex properties
 - $F(\text{String vertexLabel}, \text{long sequencedNumber}, \text{String propertyName}, \text{PropertyType type})$
 - Deterministic property value for filtered traversals.

Graph Gen: Define Graph Data Distribution

- Vertex

vertices:

- type: users.User
weight: 13
- type: trips.Trip
weight: 8
- type: documents.Document
weight: 13

- Property

properties:


- type: users.driverStatus
values:
 - value: "Active"
weight: 90
 - value: "Rejected"
weight: 5
 - value: "Waitlisted"
weight: 5

- Edge


edges:

- type: trips.droveFor
out:
 - existenceProbability: 0.1
logMean: 2.0
logSD: 0.3
- in:
 - existenceProbability: 0.9

How to deal with complicated distributions?



How to deal with complicated distributions?



Use edge distribution CSV file

degree	count
1	10000000
2	1000000
3	100000
4	30000
5	10000
6	3000
7	1000
8	600
9	200
...	
1000	1
1300	1
1800	1
20025	1
60123	1

Generate Edges

- Create edges:
 - Pick a source vertex
 - Pick a target vertex
 - Connect them
- Function of (random seed, edge distribution)

Graph Gen: Consistent Partitioned

- Need to support massive graphs.
- Support Spark to partition the generated graph.
- Each partition has its own seed.
 - $\text{partition seed} = \text{base_seed} + \text{partition_index}$
- Generate a graph of 1 billion vertex and edges in one hour.

Implemented Storages

- Graph DB with Gremlin support
- Key-value store and prefix-key-value store
 - Vertex:
 - Row key = (VertexLabel, VertexId)
 - Value: properties

Edge Storage

Key-value

- AdjacencyList
- Row Key = (EdgeLabel, sourceVertexId)
- Value : List of (targetVertexId, edge properties)

Prefix-key-value

- Multi-rows of edges
- Row key = {(EdgeLabel, **sourceVertexId**), targetVertexId}
- Value: edge properties
- To query a source vertex edges, scan rows with prefix (**EdgeLabel, sourceVertexId**).

Support Anchored Queries

- Starting from a randomly chosen vertex.
- Implemented K-hop query for key value stores.
- Or any anchored path query

startVertexLabel: users.User

queryType: gremlin

queryText: >

```
g.V(x).outE('trips.requested').subgraph('s').outV()  
.outE('trips.droveFor').subgraph('s').outV()  
.outE('documents.usedDocument').subgraph('s').outV()  
.cap('s').next()
```

Example Output

```
{
  "write.vertex": {
    "Operations": 999998,
    "Average(us)": 513.58802002204,
    "Variance(us)": 1343.5852131123756,
    "Min(us)": 52.967,
    "Max(us)": 157801.328,
    "95thPercentile(us)": 1347,
    "99thPercentile(ms)": 2
  },
  "write.edge": {
    "Operations": 628233,
    "Average(us)": 1066.3619543274551,
    "Variance(us)": 792.03133339263,
    "Min(us)": 129.731,
    "Max(us)": 33993.35,
    "95thPercentile(ms)": 2,
    "99thPercentile(ms)": 4
  },
  "read.vertex": {
    "Operations": 38349,
    "Average(us)": 1734.14784575869,
    "Variance(us)": 2459.3334678627807,
    "Min(us)": 70.626,
    "Max(us)": 58464.943,
    "95thPercentile(ms)": 4,
    "99thPercentile(ms)": 11
  },
  "read.edge": {
    "Operations": 138349,
    "Average(us)": 1573.6572573997644,
    "Variance(us)": 2037.7534999789132,
    "Min(us)": 41.721,
    "Max(us)": 56861.816,
    "95thPercentile(ms)": 4,
    "99thPercentile(ms)": 10
  },
  "subgraph": {
    "Operations": 100000,
    "Average(us)": 12439.45511393,
    "Variance(us)": 5797.011929088904,
    "Min(us)": 668.139,
    "Max(us)": 105834.563,
    "95thPercentile(ms)": 32
  },
  "subgraph.vertex.count": 67236,
  "subgraph.edge.count": 60742,
  "non.empty.subgraph.count": 6494
}
```

Open source

<https://github.com/uber/uber-graph-benchmark>

Uber Graph Benchmark

- Graph Generation
 - Schema-aware
 - Flexible distribution
 - Deterministic
 - Highly scalable thanks to Spark support
- Plugins for graph storage backend
 - Extensible for different storage systems
- Graph query
 - K-hop subgraph queries

Contributions welcome

- Test your graph database solutions
- Share your data schema, distributions and queries
- Add adapters for key-value stores
- Add adapters for graph databases
- Other:
 - Different edge generation algorithms
 - More vertex property types

Link

<https://github.com/uber/uber-graph-benchmark>

Email

- [Chris Lu](mailto:chris.lu@uber.com) <chris.lu@uber.com>
- [Joshua Shinavier](mailto:joshsh@uber.com) <joshsh@uber.com>

